



```
In [ ]: import numpy as np
import gymnasium as gym # ✓ use Gymnasium

# -----
# Create Taxi environment
# -----
env = gym.make("Taxi-v3")

print("Number of states:", env.observation_space.n) # 500
print("Number of actions:", env.action_space.n)    # 6

# Parameters
gamma = 0.9
theta = 1e-6

# -----
# Value Iteration
# -----
def value_iteration(env, gamma=0.9, theta=1e-6):
    value_table = np.zeros(env.observation_space.n)

    while True:
        delta = 0
        for state in range(env.observation_space.n):
            action_values = np.zeros(env.action_space.n)
            for action in range(env.action_space.n):
                for prob, next_state, reward, done in env.unwrapped.P[state][action]:
                    action_values[action] += prob * (reward + gamma * value_table[next_state])
            best_action_value = np.max(action_values)
            delta = max(delta, abs(value_table[state] - best_action_value))
            value_table[state] = best_action_value
        if delta < theta:
            break

    # Derive policy
    policy = np.zeros(env.observation_space.n, dtype=int)
    for state in range(env.observation_space.n):
        action_values = np.zeros(env.action_space.n)
        for action in range(env.action_space.n):
            for prob, next_state, reward, done in env.unwrapped.P[state][action]:
                action_values[action] += prob * (reward + gamma * value_table[next_state])
        policy[state] = np.argmax(action_values)

    return policy, value_table

# -----
# Policy Evaluation
# -----
def policy_evaluation(policy, env, gamma=0.9, theta=1e-6):
    value_table = np.zeros(env.observation_space.n)
    while True:
        delta = 0
        for state in range(env.observation_space.n):
```

```

        v = 0
        action = policy[state]
        for prob, next_state, reward, done in env.unwrapped.P[state][action]:
            v += prob * (reward + gamma * value_table[next_state])
        delta = max(delta, abs(value_table[state] - v))
        value_table[state] = v
    if delta < theta:
        break
    return value_table

# -----
# Policy Improvement
# -----
def policy_improvement(value_table, policy, env, gamma=0.9):
    policy_stable = True
    for state in range(env.observation_space.n):
        old_action = policy[state]
        action_values = np.zeros(env.action_space.n)
        for action in range(env.action_space.n):
            for prob, next_state, reward, done in env.unwrapped.P[state][action]:
                action_values[action] += prob * (reward + gamma * value_table[next_state])
        policy[state] = np.argmax(action_values)
        if old_action != policy[state]:
            policy_stable = False
    return policy, policy_stable

# -----
# Policy Iteration
# -----
def policy_iteration(env, gamma=0.9, theta=1e-6):
    policy = np.random.choice(env.action_space.n, size=env.observation_space.n)
    value_table = np.zeros(env.observation_space.n)

    while True:
        value_table = policy_evaluation(policy, env, gamma, theta)
        policy, policy_stable = policy_improvement(value_table, policy, env, gamma)
        if policy_stable:
            return policy, value_table

# -----
# Run Both Algorithms
# -----
print("\n Running Policy Iteration...")
pi_policy, pi_value = policy_iteration(env, gamma, theta)
print("Policy Iteration: Optimal Value Function shape =", pi_value.shape)
print("Policy Iteration: Optimal Policy shape =", pi_policy.shape)

print("\n Running Value Iteration...")
vi_policy, vi_value = value_iteration(env, gamma, theta)
print("Value Iteration: Optimal Value Function shape =", vi_value.shape)
print("Value Iteration: Optimal Policy shape =", vi_policy.shape)

# -----

```

```

# Quick check: same optimal policy?
# -----
print("\n Do both methods give same optimal policy? ->", np.array_equal(pi_pol

# -----
# Demonstration: play one episode using optimal policy
# -----
state, _ = env.reset()
done = False
total_reward = 0
steps = 0

print("\n Demo run with Optimal Policy:")
while not done and steps < 20:
    action = pi_policy[state]
    state, reward, terminated, truncated, _ = env.step(action)
    done = terminated or truncated
    total_reward += reward
    steps += 1
    env.render()

print("Episode finished in", steps, "steps with reward:", total_reward)

```

Number of states: 500

Number of actions: 6

Running Policy Iteration...

Policy Iteration: Optimal Value Function shape = (500,)

Policy Iteration: Optimal Policy shape = (500,)

Running Value Iteration...

Value Iteration: Optimal Value Function shape = (500,)

Value Iteration: Optimal Policy shape = (500,)

Do both methods give same optimal policy? -> True

Demo run with Optimal Policy:

Episode finished in 12 steps with reward: 9

```

/usr/local/lib/python3.12/dist-packages/gymnasium/envs/toy_text/taxi.py:443: Us
erWarning: WARN: You are calling render method without specifying any render mo
de. You can specify the render_mode at initialization, e.g. gym.make("Taxi-v3",
render_mode="rgb_array")
  gym.logger.warn(

```